

Hojin Park

hojinp.github.io

Email: hojinp@andrew.cmu.edu

Mobile: +1-617-642-1884

RESEARCH INTEREST

Distributed systems, Public cloud storage, Storage systems, Cache, Deep learning systems

EDUCATION

| | |
|---|--------------------------|
| Carnegie Mellon University (CMU) | Sep. 2019 – Present. |
| Ph.D. Student in Computer Science Department | Pittsburgh, PA |
| Seoul National University (SNU) | Mar. 2013 – Feb. 2019 |
| B.S. in Electrical and Computer Engineering | Seoul, Republic of Korea |
| Graduated with Summa Cum Laude (GPA: 4.21 / 4.30) | |
| Korea Science Academy of KAIST | Mar. 2010 – Mar. 2013 |
| Math & Science specialized high school | Busan, Republic of Korea |

RESEARCH EXPERIENCE

| | |
|--|-----------------------|
| Parallel Data Lab, CMU | Sept. 2019 - Present. |
| Graduate Research Assistant (Advisor: Prof. George Amvrosiadis, Prof. Greg Ganger) | Pittsburgh, PA |

- **Macaron:** multi-cloud/region aware cache auto-configuration system
 - I am developing a system that optimizes the cache configuration in runtime to minimize the total cost of accessing data in a remote cloud or region.
- **Mimir:** a tool that finds the cost-efficient cloud storage configuration for storage systems
 - I developed Mimir, a system that helps users to make optimal decisions when composing distributed storage systems in the public cloud.
 - Mimir lets users enter a set of SLOs and outputs the most cost-efficient cloud resources configuration that minimizes the overall cost paid by the user.
- **Burstable storage in public clouds:** proposes a new way of exploiting burstable storage service cost-efficiently
 - I examined how burstable storage can be leveraged to reduce cost and/or improve performance for three use cases with different data-longevity requirements: traditional persistent storage, caching, and ephemeral storage
 - I found that by aggressively exploiting burstable storage service in public clouds, it is possible to increase storage throughput by up to 100x at a cost increase of only 10-40%.

| | |
|---|--------------|
| VMware Research | Summer, 2023 |
| Intern - Member of Technical Staff (Mentor: Adriana Szekeres) | Remote |

- **Trace analysis:** Generated object storage access traces for cross-cloud/region caching research and analyzed data access patterns

| | |
|---|--------------|
| CORE group, Microsoft Research | Summer, 2021 |
| Research Intern (Mentor: Ishai Menache) | Virtual |

- **VM packing:** Researched a deep learning model that optimizes the VM packing efficiency in a large cluster

| | |
|---|--------------------------|
| Software Platform Laboratory, SNU | Jan. 2017 – Aug. 2019 |
| Research Intern (Advisor: Prof. Byung-Gon Chun) | Seoul, Republic of Korea |

- **Parallax:** a tool for automatic parallelization of deep learning training
 - Transforms a single-GPU deep learning model for distributed execution, handling correctness and scalability.
 - I used Parallax to explore two distributed training designs: Parameter Server and AllReduce.
 - I implemented four deep learning models with each distributed architecture to better understand these designs.
 - I ran experiments to evaluate these models on Parallax, in terms of correctness, scalability, and optimization.
- **Cruise:** a distributed machine learning framework with automatic system configuration

- Optimizes a system by adjusting worker/server assignment to homogeneous clusters at runtime.
- I enabled Cruise to work with heterogeneous cluster sets by implementing a custom linear-programming-based solver to optimize a generalized cost model.
- I implemented a Gradient Boosting Tree (GBT) application on top of Cruise.

Virtual Machine and Optimization Laboratory, SNU
Research Intern (Advisor: Prof. Soo-Mook Moon)

Jan. 2018 – July. 2018
 Seoul, Republic of Korea

- **GitChain:** a distributed version control system using blockchain
 - Uses a public ledger to save version controlled repositories in InterPlanetary File System (IPFS).
 - I designed and implemented blockchain-related components of the system.
 - I implemented basic Git functions, such as push, pull, and clone, on the IPFS.

SCHOLARSHIPS & AWARDS

| | |
|--|---|
| International Graduate Student Scholarship <i>Full tuition, insurance, and living expenses (5 years)</i> | Sep. 2019 - Aug. 2024 <i>Korea Foundation for Advanced Studies</i> |
| Blockchain Technology Competition <i>Two-person team won first prize (\$3,000), with GitChain project</i> | Jul. 2018 <i>LINE, KIISE</i> |
| Undergraduate Study Scholarship <i>Full tuition and stipend (\$2,500/semester)</i> | Feb. 2017 - Dec. 2018 <i>Kwanjeong Educational Foundation</i> |
| Academic Excellence Scholarship <i>Full tuition</i> | Jun. 2013 - Dec. 2014 <i>SNU</i> |

PROGRAMMING SKILLS

- Languages: C/C++, Java, Python, SQL
- Multicore/GPU Libraries: OpenCL, CUDA, MPI, OpenMP
- Other: Tensorflow, Horovod, Gurobi (ILP)

PUBLICATIONS AND PREPRINTS

- [1] Hojin Park, Ziyue Qiu, Gregory R. Ganger, George Amvrosiadis. Reducing cross-cloud/region costs with the auto-configuring MACARON cache. *SOSP 2024*, November 2024.
- [2] Hojin Park, Gregory R. Ganger, George Amvrosiadis. Mimir: Finding Cost-efficient Storage Configurations in the Public Cloud. *SYSTOR 2023*, June 2023.
- [3] Hojin Park, Gregory R. Ganger, George Amvrosiadis. More IOPS for Less: Exploiting Burstable Storage in Public Clouds. *HotCloud 2020*, July 2020.
- [4] Woo-Yeon Lee, Yunseong Lee, Joo Seong Jeong, Gyeong-In Yu, Joo Yeon Kim, Hojin Park, Beomyeol Jeon, Wonwook Song, Gunhee Kim, Markus Weimer, Brian Cho, Byung-Gon Chun. Automating System Configuration of Distributed Machine Learning. *ICDCS 2019*, March 2019.
- [5] Soojeong Kim, Gyeong-In Yu, Hojin Park, Sungwoo Cho, Eunji Jeong, Hyeonmin Ha, Sanha Lee, Joo Seong Jeong, Byung-Gon Chun. Parallax: Sparsity-aware Data Parallel Training of Deep Neural Networks. *EuroSys' 19*, March 2019.
- [6] Soojeong Kim, Eunji Jeong, Joo Seong Jeong, Gyeong-In Yu, Hojin Park, Byung-Gon Chun. Auto-Parallelizing Deep Learning for Multi-machine, Multi-GPU Environments. *Workshop on AI Systems at Symposium on Operating Systems Principles (SOSP)*, October 2017.

TEACHING

Carnegie Mellon University
 Storage Systems (15-746)
 Advanced Cloud Computing (15-719)

Fall 2022
 Spring 2022

TALKS

MACARON: Multi-cloud/region Aware Cache Auto-Reconfiguration

- The 30th ACM Symposium on Operating Systems Principles (SOSP) November 2024
- CMU Parallel Data Lab Retreat October 2024
- CMU Advanced Cloud Computing Course Guest Lecture April 2024
- Alluxio & Uber Data Infra Meetup January 2024
- CMU Parallel Data Lab Retreat November 2023

Toward cost-efficient storage systems and data transfer in public clouds

- Salesforce Database Team Reading Group January 2024

Mimir: Finding Cost-efficient Storage Configurations in the Public Cloud

- The 16th ACM International Systems and Storage Conference (SYSTOR) June 2023
- CMU Advanced Cloud Computing Course Guest Lecture April 2023
- CMU Parallel Data Lab Retreat November 2022

MENTORING

- Saileshwar Karthik (CMU Information Networking Institute masters student) 2025
- Mohit Gaggar (CMU Information Networking Institute masters student) 2025
- Fulun Ma (CMU Computational Data Science masters student) 2024
- Somansh Satish (CMU Computational Data Science masters student) 2023
- Anurag Choudhary (CMU Computational Data Science masters student) 2023
- Midhush Manohar Thevendria Karthic (CMU Computational Data Science masters student) 2023
- Shalini Shukla (CMU ECE masters student) 2022
- Hao Yang Lu (CMU SCS masters student) 2022